

CAAM 520: COMPUTATIONAL SCIENCE II

HOMEWORK 2.

WEI WU

1. INTRODUCTION

In this project, we build upon the previous project: we parallelize our linear system solver using MPI.

1.1. Mesh Grid Partition. We partition the mesh grid by the rows, illustrated by the picture below.




FIGURE 1. Mesh grid partition

However, I made the design decision to only partition the interior nodes. Let N be the number of interior nodes in a single row, and let s be the number of processors used. I try to divide the interior nodes by rows evenly among all processors. If rm is the remainder of N divided by s , I allocate extra rm rows to processor 0. Two extra rows of ghost nodes, one at the bottom and the other at the top, are allocated for the local u .

```
/*
Design decision: allocate extra N%size rows to processor 0,
and N/size + 2 to all other processes.
*/
int local_row_size = N/size + 2;
if (rank == 0){
    local_row_size += N%size;
}

double *local_u = (double*) calloc((local_row_size)*(N+2), sizeof(double));
double *f = (double*) calloc((local_row_size)*(N+2), sizeof(double));
```

```

double h = 2.0/(N+1);

for (int i = 0; i < N+2; ++i){
    for (int j = 0; j <
        local_row_size; ++j){
        const double x = -1.0 + i*h;
        double y = -1.0 + j*h;
        if (rank > 0)
            y += (rank - 1)*(N/size)*h + (N/size + N%size)*h;
        f[i + j*(N+2)] = sin(PI*x)*sin(PI*y) * h*h;
    }
}

```

1.2. **Parallel weighted Jacobi.** A very high level pseudocode of the parallel weighted Jacobi:

Algorithm 1 My algorithm

```

1: procedure MYALGO( $N, tol, size, rank$ )
2:   while  $global\_res \geq tol$  do
3:     Send up unless  $rank = size - 1$ , then receive from  $rank - 1$ 
4:     Send down unless  $rank = 0$ , then receive from  $rank + 1$ 
5:     Compute  $u\_local$  using weighted Jacobi; compute  $local\_res$ 
6:     Compute  $global\_res$ 
7:   return  $u\_local$ 

```

The above pseudocode corresponds to the following code snippet:

```

/*
Ghost nodes are stored on
u[0,...,N+1]
and
u[(local_row_size-1)*(N+2),...,(local_row_size)*(N+2)-1]
Note that for the bottom and top, the exterior nodes will be
ghost nodes.
*/
double *unew = (double*)calloc((N+2)*(local_row_size),sizeof(double));

double w = 1.0;
double invD = 1./4.; // factor of h cancels out

double global_res2 = 1.0;
unsigned int iter = 0;
MPI_Status status;
while(global_res2>tol*tol){

    double local_res2 = 0.0;
    for (int i=0; i<= N+1; ++i){
        /*Send up unless at the top, then receive from below*/
        if (rank < size - 1)

```

```

MPI_Send(&u[i + (local_row_size-2)*(N+2)], 1, MPI_DOUBLE,
        rank + 1, 0, MPI_COMM_WORLD);
if (rank > 0)
    MPI_Recv(&u[i], 1, MPI_DOUBLE, rank-1, 0, MPI_COMM_WORLD, &status);
/*Send down unless at the bottom, then receive from above*/
if (rank > 0)
    MPI_Send(&u[i + 1*(N+2)], 1, MPI_DOUBLE,
            rank-1, 1, MPI_COMM_WORLD);
if (rank < size - 1)
    MPI_Recv(&u[i + (local_row_size-1)*(N+2)], 1, MPI_DOUBLE,
            rank + 1, 1, MPI_COMM_WORLD, &status);
}

// update interior nodes using Jacobi; does not touch ghost nodes
// int global_id;
for(int i=1; i<=N; ++i){
    for(int j=1; j<=local_row_size-2; ++j){
        const int id = i + j*(N+2); // x-index first
        const double Ru = -u[id-(N+2)]-u[id+(N+2)]-u[id-1]-u[id+1];
        const double rhs = invD*(f[id]-Ru);
        const double oldu = u[id];
        const double newu = w*rhs + (1.0-w)*oldu;
        local_res2 += (newu-oldu)*(newu-oldu);
        unew[id] = newu;
    }
}

for (int i = 0; i < (N+2)*(local_row_size); ++i){
    u[i] = unew[i];
}

++iter;

// calcualte global residual
MPI_Allreduce(&local_res2, &global_res2, 1, MPI_DOUBLE, MPI_SUM,
MPI_COMM_WORLD);

```

1.3. Debugging. I did not use dbg to debug. I extensively used printf and return 0 to single out chunks of code. I checked correctness of the program on a small problem with $N = 2$, just as last time.

2. RESULTS ON XPS 13 9370

For a residual $< 1e - 6$, $N = 10$. The max error and number of iterations stays the same. Depending on the number of processors, the total time could be either higher or lower than using a single processor. See tables belows for some of the results.

TABLE 1. $N = 10$, $\text{tol} = 1\text{e-}6$

Num. Processors	Iter	Max Error	Time(s)
1	64	0.00137	0.00137
2	64	0.00137	0.000769138
3	64	0.00137	0.003733
4	64	0.00137	0.000908077
5	64	0.00137	0.00152731
6	64	0.00137	0.00202286
7	64	0.00137	0.00161999
8	64	0.00137	0.14174

TABLE 2. $N = 100$, $\text{tol} = 1\text{e-}6$

Num. Processors	Iter	Max Error	Time(s)
1	4395	6.13072e-06	0.88749
2	4395	6.13072e-06	0.635089
3	4395	6.13072e-06	0.650189
4	4395	6.13072e-06	0.567912
5	4395	6.13072e-06	0.609207
6	4395	6.13072e-06	0.589206
7	4395	6.13072e-06	0.580673
8	4395	6.13072e-06	0.650094