

CAAM 520 · Computational Science II

Spring 2019 · Rice University

[MAIN PAGE](#) // [HOMEWORK](#)

Notes, Codes, and Supplementary Material

Lecture OCCA lab: install OCCA on both Ubuntu and NOTS.

33 OCCA 2-step installation [instructions](#)

(4/17):

Installation on NOTS:

- get a GPU using the interactive shell (--reservation=CAAM520_2)
- Run "module load CUDA"
- Follow the installation procedure [here](#)
- Add "export
LD_LIBRARY_PATH="\${LD_LIBRARY_PATH}:/home/your_username/occa/lib""
to your "~/.bashrc" file. This will ensure OCCA remains in your path when you log back in
- "cd examples/cpp/1_add_vectors/", type "make -j", and run ./main to est.

Installation on Ubuntu:

- Follow the installation procedure [here](#)
- Add "export
LD_LIBRARY_PATH="\${LD_LIBRARY_PATH}:/home/your_username/occa/lib""
to your "~/.bashrc" file. This will ensure OCCA remains in your path when you log back in
- "cd examples/cpp/1_add_vectors/", type "make -j", and run ./main to est.
- You may observe that the examples run, but return a segmentation fault. This is likely due to a [documented issue](#) with Intel's OpenCL installation + the OpenCL command clReleaseProgram (which is run during clean-up by OCCA)

[okl_demo.zip](#): example showing the use of inner/outer loops, shared and exclusive memory in OCCA

. Once you've verified the installation, you may work on the EC portion of HW 5 involving implementations of Jacobi + reduction in OCCA.

Lecture Structure and syntax of OCCA.

32 Implicit "outer" and "inner" for loops

(4/15): Implicit synchronization between loops

New memory space: exclusive variables are thread-local but persist between loops.

Lecture Finishing up OpenCL, introduction to OCCA.

31 Differences between OCCA and OpenCL: translation into native languages for
(4/12): OCCA
OCCA website

Lecture OpenCL lab.

30 Goal: install OpenCL on your Ubuntu VirtualBox installation, test with an OpenCL
(4/10): implementation of reduce.cu

Main steps (based off of this [askUbuntu answer](#)):

- Download the [OpenCL SDK](#) for Intel CPUs (if you have another type of CPU, you may need a different SDK).

- Install OpenCL libraries and tools using

- `sudo apt install ocl-icd-libopencl1`
- `sudo apt install opencl-headers`
- `sudo apt install clinfo`
- `sudo apt install ocl-icd-opencl-dev`

- Install package to convert .rpm to .deb files:

- `sudo apt-get install -y rpm alien libnuma1`

- Untar downloaded OpenCL SDK files

- `tar -xvf opencl_runtime_16.1.1_x64_ubuntu_6.4.0.25.tgz`

- Turn rpm files to deb

- `cd opencl_runtime_16.1.1_x64_ubuntu_6.4.0.25/rpm/`
- `fakeroot alien --to-deb opencl-1.2-base-6.4.0.25-1.x86_64.rpm`
- `fakeroot alien --to-deb opencl-1.2-intel-cpu-6.4.0.25-1.x86_64.rpm`

- Install .deb packages

- `sudo dpkg -i opencl-1.2-base_6.4.0.25-2_amd64.deb`
- `sudo dpkg -i opencl-1.2-intel-cpu_6.4.0.25-2_amd64.deb`

- Create local config file

- `sudo touch /etc/ld.so.conf.d/intelOpenCL.conf`

- Open the file and add OpenCL config info

- `sudo emacs -nw /etc/ld.so.conf.d/intelOpenCL.conf`
- Type `"/opt/intel/opencl-1.2-6.4.0.25/lib64/clinfo"` in the file and close

- Create a "vendors" dir and create an icd link

- `sudo mkdir -p /etc/OpenCL/vendors`
- `sudo ln /opt/intel/opencl-1.2-6.4.0.25/etc/intel64.icd /etc/OpenCL/vendors/intel64.icd`
- `sudo ldconfig`

- Test the installation:

- Build `"cldevices.cpp"` using `"g++ cldevices.cpp -o cldevices -lOpenCL"`

- Run `./cldevices` to view available devices (should just be a CPU)

End goal of the lab: convert the optimized `reduce.cu` CUDA kernel to working OpenCL code. Test the code on Ubuntu (VirtualBox), and if possible, on NOTS.

Lecture Converting CUDA to OpenCL, timing OpenCL code

29 `clmxm_timing.cpp`: OpenCL matrix-matrix multiplication code with OpenCL event timing.
(4/8):

`mxm.cl`: kernel file for above example. Converting between CUDA and OpenCL: a non-comprehensive list of changes

- `"__global__"` keyword for kernels becomes `"__kernel"`
- Add `"__global"` to pointer arguments to global (DRAM) memory
- `"__shared__"` keyword for shared memory becomes `"__local"`
- `"__syncthreads()"` becomes `"barrier(CLK_LOCAL_MEM_FENCE)"`

To build and run OpenCL code on NOTS, use `"g++ -I$EBROOTCUDA/include clmxm_timing.cpp -lOpenCL"`. This will use NVIDIA's OpenCL implementation

Lecture More on OpenCL Wrappers to simplify OpenCL host code (setup and kernel build functions)

(4/5): `cldevices.cpp`: code to determine OpenCL devices available
`cldemo.cpp`: (un-simplified) code demonstrating building and running of an OpenCL kernel

`foo.cl`: kernel for above code. `cldemo.cpp`: simplified code demonstrating building and running of an OpenCL transpose kernel

`transpose.cl`: kernel for above code.

Lecture Introduction to OpenCL (Chapter 2 of OpenCL book)

27 [OpenCL book pdf](#).

(3/29):

Steps to an OpenCL program:

Setup

- choose platform and device (`clGet*Info`, `clGetPlatformIDs`, `clGetDeviceIDs`)
- combine platform and device to create compute environment (`clCreateContext`)
- create "stream" (`clCreateCommandQueue`).
- Options for sequential or out-of-order scheduling (`CL_QUEUE_PROFILING_ENABLE` / `CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE`).

Build program

- create program object from one or more kernel sources (`clCreateProgramWithSource`)
- program object gets built into executables (`clBuildProgram`).
- build executable kernel from program (`clCreateKernel`)

Allocate memory

- allocate memory (`clCreateBuffer`). Specify type of memory
- `CL_MEM_COPY_HOST_PTR` = memcpy from specified host array
- `CL_MEM_USE_HOST_PTR` = specify host pointer. In some implementations, uses pinned memory.
- `CL_MEM_ALLOC_HOST_PTR` = pinned memory (host-accessible)

Queue kernels, set arguments

- set kernel arguments (`clSetKernelArg`)
- queue kernel to run (`clEnqueueNDRangeKernel`).
- get result back (`clEnqueueReadBuffer`)
- not used in demo: write buffer (`clEnqueueWriteBuffer`)

Synchronize and clean up

- `clFlush` - blocks until commands in queue have executed (not completed!)
- `clFinish` - blocks until commands in queue have finished
- Context creates queue, program, memory buffers. Program creates kernels. Kernels + memory feed into queue.

Lecture CPU/GPU data transfers

- 26 [pinnedMemoryExample.cu](#): example of faster transfers using pinned memory and
(3/27): `cudaMallocHost`
Can overlap CPU computation with CUDA kernels trivially
Multiple CUDA streams + pinned memory allow overlap of GPU compute/memory transfer
[async.cu](#): example from the CUDA samples on overlapping GPU computation with data transfer using 4 CUDA streams.

Lecture Optimizing matrix multiplication with shared memory

- 25 Matrix multiplication with shared memory tiling.
(3/25): [mxm.cu](#): code with three versions of matrix transposition.
Discussion of full CUBLAS strategy for optimizing matrix-matrix multiplication

Lecture Optimizing matrix transposes with shared memory

- 24 Bank conflicts for 32-by-32 arrays
(3/22): [transpose.cu](#): code with three versions of matrix transposition (global memory, shared memory, shared memory with bank conflict treatment)

Lecture Even more on profiling GPU kernels for efficiency

- 23 Roofline model: compute bound vs memory bound kernels
(3/20): Peak vs realistic performance
Comparison of matrix-matrix multiplication with CUBLAS
[matmult_cublas.cu](#): example code for CUBLAS (compilation instructions in header)
[ilp.cu](#): optimizing code for adding two vectors by increasing work per thread.
[roofline.m](#): Matlab code for computing roofline plots.

Lecture More on profiling GPU kernels for efficiency

- 22 Roofline model: compute bound vs memory bound kernels
(3/18): Peak vs realistic performance

Lecture Profiling GPU kernels for efficiency

- 21 GPU occupancy: [online occupancy calculator](#)
(3/8): `nvprof` (Nvidia profiler) for computing timings, number of floating point operations, bandwidth.

Lecture Optimizing a reduction kernel

- 20 Concepts: warp divergence, shared memory bank conflicts.
(3/6): Example optimized reduction code: [reduce.cu](#)

Reduction code based off of Mark Harris' optimized reduction kernel talk.

Lecture More on shared memory

19 Matvec example code: [matvec.cu](#)

(3/4): Stencil example code: [stencil.cu](#)

Lecture Introduction to shared memory

18 [matvec.cu](#)

(2/21):

Lecture More on GPU computing

17 Matrix-matrix multiplication: [matmult.cu](#)

(2/19): Global GPU memory and coalesced memory access

Introduction to nvprof (Nvidia profiler) for timing

Kirk and Hwu Chapter 3.

Lecture Introduction to GPU computing

16 Host/device setup

(2/18): Nvidia, AMD GPUs: CUDA vs OpenCL

[add.cu](#): CUDA code to add two vectors. Compile using "nvcc add.cu -o add", run as usual ("./add").

To run on NOTS in interactive mode: "srun --pty --partition=commons --gres=gpu:1 --time=1:00:00 \$SHELL" will request a single GPU for 1 hour.

Lecture Divide and conquer using OpenMP tasks

15 Recursive blocked matrix-matrix multiplication [code](#)

(2/18):

Lecture More advanced OpenMP: task-based parallelism

14 OpenMP tutorial with several examples

(2/15): [An extensive list of more OpenMP resources](#)

Lecture Code optimization

13 [matmat_omp.c](#): inefficient matrix-matrix multiplication code

(2/13): [dot_omp.c](#): optimized matrix-matrix multiplication code

Cache memory, data layouts, compiler optimization flags

Assorted OpenMP options: nested parallelism, loop collapsing (e.g. `collapse(2)`), conditional statements in OMP pragmas.

Lecture Introduction to shared memory parallelism and OpenMP

12 [helloworld_omp.c](#): hello world using OpenMP

(2/11): [dot_omp.c](#): computing a dot product

Enabling OpenMP: "gcc -fopenmp ..."

Pragmas, parallel for loops, shared and private variables, race conditions.

Lecture Domain decomposition using graph partitioning

11 Link to [METIS](#) by the Karypis lab.

(2/6): Installation instructions: install Cmake (sudo apt-get install cmake), cd into metis-5.1.0/, type "make config" and "make".

[metisDriver.c](#): example code using Metis to partition a structured grid.

Other methods for load-balancing and domain decomposition: [space-filling curves](#) (see Figures 1,2 for illustrations)

Lecture Debugging MPI

10 Debug option 1: launch multiple gdb processes using "mpiexec -n 2 xterm -e gdb

(2/4): `./a.out`
Debug option 2: automatically attach gdb (see `debug_mpi.c`, run using `"mpiexec -n 2 ./a.out"`)
Open-MPI FAQ for debugging MPI.

Lecture MPI_Probe, MPI_Getcount, timing

9 `probe.c`: demo of MPI_Probe, MPI_Getcount

(1/31): `matvec2.c` for timing and `matvec_timing.slurm` for running on NOTS (must compile `matvec2.c` first).

Lecture Parallel even-odd sort

8 Even-odd sort code based on Pachecho 3.7.2: `parsort.c`.

(1/25):

Lecture Lab on matrix-vector products, NOTS

7 Description of the lab.

(1/23): Matrix-vector product code `matvec.c` using row-based storage.

Lecture More on collective communication in MPI

6 Butterfly communication: MPI_Allreduce, MPI_Allgather, MPI_Alltoall

(1/18): `all_collectives.c`: demo of above routines.

Parallel matrix-vector products

Using NOTS

Lecture Collective communication in MPI

5 MPI_Reduce, MPI_Bcast, MPI_Barrier, MPI_Gather, MPI_Scatter

(1/16): `collectives.c`: demo of above routines.

Butterfly communication: MPI_Allreduce, MPI_Allgather, MPI_Alltoall

Lecture Hw 1, reductions in MPI.

4 `MPI_REDUCE reduce.c`: computes a reduction using tree parallelism and compares

(1/14): the result to MPI_REDUCE.

Lecture Interactive lab.

3 Set up VirtualBox (<https://www.virtualbox.org/>).

(1/11): - Set up Ubuntu 16.04 on VirtualBox (64 bit at <http://releases.ubuntu.com/16.04/ubuntu-16.04.5-desktop-amd64.iso>, 32 bit at <http://releases.ubuntu.com/16.04/ubuntu-16.04.5-desktop-i386.iso>)
- Information for running 64 bit on a 32 bit machine from Seth Brown (optional): [Link](#)
- I recommend setting up VM with at least half system memory for performance.
- Install emacs: `sudo apt install emacs` (or VIM).
- Install MPICH2: `sudo apt-get install mpich` (should be MPICH2)
- Install Git: `sudo apt-get install git`

Set up Github repository at <https://github.com/>

- Create git repo on Github

- Check out Git repo in VirtualBox (`cd ~, git clone https://github.com/*github_id/*repo_address*`)

- Edit your README file to add your name, then commit (`git add README.md, git commit -m "committing update to README file", git push`)

- Try running MPI demo programs. Given time, write "ping-pong" example.

- add "jlchan" and "ctl" as Collaborators to your Github repository.

Lab: construct a "ping-pong" MPI program with 2 ranks which sends a ping_pong

variable between each rank, incrementing it until ping_pong is larger than 10.

Lecture More message passing with MPI.

2 (1/9): MPI_Send/Recv behavior, MPI_Sendrecv

Issues with message passing (deadlock)

Nonblocking MPI communication: MPI_Isend, MPI_Irecv, MPI_Wait

[even_odd.c](#)

Pacheco 3.3

Lecture Introduction to MPI (Chapter 3 in Pacheco)

1 (1/7): Simple MPI programs, [myhelloworld.c](#)

Initializing an MPI program, MPI_Send, MPI_Recv.

Piazza forum [signup](#).

Pacheco 3.1